

Model Evaluation & Choosing the Right Model

Why Evaluation Matters

Training a model is easy. Knowing if it actually works is the hard part. A model that scores 98% on training data might score 62% on new data. This gap is the real test of whether your model has learned a general pattern or just memorised the training examples.

Building a Machine Learning model is only half the job.

The real question is:

"Will the model work well on new, unseen data?"

A model may achieve:

98% Training Accuracy 62% Test Accuracy

This means the model has memorized the training data instead of learning general patterns.

Therefore, evaluating a model correctly is essential before deploying it in the real world.

Types of Error

Error Type	What It Measures	What It Tells You
Training Error	Error on training data	Low error means the model learned the training data (or mer
Validation Error	Error on validation data	Helps detect overfitting and tune the model.
Test Error	Error on completely unseen data	Best estimate of real-world performance.

Concept	Key Idea
Training Error	Measures performance on seen data
Validation Error	Helps tune and improve models
Test Error	Estimates real-world performance
Overfitting	Excellent training performance but poor test performance
Golden Rule	Never repeatedly use the test set for model tuning

The Golden Rule

NEVER use your test set more than once. Every time you peek at test results and adjust your model, you're leaking information. The final test score must come from data the model has NEVER influenced in any way

See Training Error vs Test Error Gap Watch the gap grow as you increase model complexity:

```
# ■ Training Error vs Test Error Gap ■
from sklearn.datasets import make_classification
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt, numpy as np

X, y = make_classification(
    n_samples=500,
    n_features=10,
    random_state=42)

X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.25,
    random_state=42)

train_errs, test_errs = [], []

depths = range(1, 25)

for d in depths:

    tree = DecisionTreeClassifier(
        max_depth=d,
        random_state=42)

    tree.fit(X_train, y_train)

    train_errs.append(1 - accuracy_score(y_train, tree.predict(X_train)))
    test_errs.append(1 - accuracy_score(y_test, tree.predict(X_test)))

best = depths[test_errs.index(min(test_errs))]


plt.figure(figsize=(9, 5))
plt.plot(depths, train_errs, 'o-', lw=2, color='steelblue', label='Training')
plt.plot(depths, test_errs, 's-', lw=2, color='tomato', label='Test Error')
plt.axvline(x=best, color='green', lw=2, linestyle='--', label=f'Best depth=')
plt.fill_between(depths, train_errs, test_errs, alpha=0.1, color='red', label='')
plt.xlabel('Tree Depth (model complexity)', fontsize=11)
```

```
plt.ylabel('Error Rate (1 - Accuracy)', fontsize=11)
plt.title('Training Error vs Test Error – The Overfitting Story', fontsize=11)
plt.legend(); plt.grid(alpha=0.3); plt.tight_layout(); plt.show()
# STORY: As depth grows, train error → 0 (perfect memorisation)
# But test error has a minimum – that's the sweet spot
```



Cross-Validation

A single train/test split gives you ONE score — but what if that split was unlucky? Maybe all the hard examples landed in your test set. K-Fold Cross-Validation solves this by using K different splits and averaging the scores.



Each fold is used as the test set exactly once. The final score is the mean of all 5 fold scores.

K Value	Trade-Off	Recommended Use
K = 5	Fast and reliable	Default choice; works well for most datasets
K = 10	More reliable but slower	Preferred for smaller datasets (< 1000 rows)

K Value	Trade-Off	Recommended Use
K = 3	Very fast but slightly less stable	Large datasets (> 100k rows) where 5-fold CV be
K = N (Leave-One-Out CV)	Most reliable but extremely slow	Very small datasets (< 100 rows) where every sam

K-Fold Cross Validation Idea

Instead of testing only once:

Divide data into K equal parts (folds)

Use one fold as the test set

Use remaining folds as training data

Repeat until every fold has been used as a test set

Average all scores

Cross-Validation Three Ways

From the simplest one-liner to full manual control:

```
# ■ Cross-Validation: One-liner + StratifiedKFold + Pipeline ■
from sklearn.datasets import make_classification
from sklearn.model_selection import (cross_val_score, KFold,
StratifiedKFold, cross_validate)
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
import numpy as np
```

```
X, y = make_classification(n_samples=600, n_features=10, random_state=42)
model = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
# ■ Method 1: Simplest one-liner ■
scores = cross_val_score(model, X, y, cv=5, scoring='accuracy')
print(f'5-Fold CV Scores: {scores.round(3)}')
print(f'Mean: {scores.mean():.3f} Std: {scores.std():.3f}')
print(f'Reliable range: {scores.mean()-2*scores.std():.3f} - {scores.mean()+
```

```
5-Fold CV Scores: [0.858 0.933 0.908 0.9 0.942]
Mean: 0.908 Std: 0.029
Reliable range: 0.850 - 0.967
```

```
# ■ Method 2: StratifiedKFold (keeps class ratio in each fold) ■
# Use this for classification, especially with imbalanced classes
skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
strat_scores = cross_val_score(model, X, y, cv=skf, scoring='f1_macro')
print(f'\nStratified 5-Fold F1: {strat_scores.mean():.3f} ± {strat_scores.st
```

Stratified 5-Fold F1: 0.908 ± 0.027

```
# ■ Method 3: cross_validate – get multiple metrics at once ■
results = cross_validate(model, X, y, cv=5,
                          scoring=['accuracy', 'f1_macro', 'roc_auc'],
                          return_train_score=True)
print('\ncross_validate results:')
for metric in ['accuracy', 'f1_macro', 'roc_auc']:
    train_k = f'train_{metric}'
    test_k = f'test_{metric}'
    print(f' {metric:12s}: train={results[train_k].mean():.3f} test={results[t
```


```
cross_validate results:
accuracy      : train=1.000 test=0.908
f1_macro      : train=1.000 test=0.908
roc_auc       : train=1.000 test=0.970
```

```
# ■ Method 4: Pipeline + CV (scale inside CV – no data leakage!) ■
pipe = Pipeline([
    ('scaler', StandardScaler()), # Scale inside each fold
    ('model', LogisticRegression(max_iter=1000))
])
pipe_scores = cross_val_score(pipe, X, y, cv=5, scoring='accuracy')
print(f'\nPipeline CV (no leakage): {pipe_scores.mean():.3f} ± {pipe_scores.
```

Pipeline CV (no leakage): 0.915 ± 0.030

Comparing Multiple Models

There is no single best algorithm for every problem. The right approach is to train several models, evaluate them on the same data with the same metric, and choose based on performance, speed, and interpretability needs.



Random Forest wins on accuracy, but Decision Tree is fastest and most interpretable.

Model	Best For	
Logistic Regression	Linearly separable data, binary classification, probability prediction	Data contain
Decision Tree	Need easy interpretation, rule-based decisions, mixed feature types	Large data
Random Forest	General-purpose classification, high accuracy, robust predictions	Very large c

Model	Best For
K-Nearest Neighbors (KNN)	Small datasets, similarity-based problems, simple classification tasks

Build a Model Leaderboard with One Loop

Train 5 models on the same data, collect metrics, build a comparison DataFrame:

```
# ■ Model Comparison Leaderboard ■
from sklearn.datasets import make_classification
from sklearn.model_selection import StratifiedKFold, cross_validate
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
import pandas as pd, numpy as np, matplotlib.pyplot as plt
X, y = make_classification(n_samples=800, n_features=15, random_state=42)
```

```
# ■ Define models (wrapped in Pipeline to include scaling) ■
models = {
    'Logistic Regression' : Pipeline([('sc',StandardScaler()),('m',LogisticRegression())],
    'Decision Tree' : DecisionTreeClassifier(max_depth=5, random_state=42),
    'Random Forest' : RandomForestClassifier(n_estimators=100, random_state=42),
    'KNN (k=5)' : Pipeline([('sc',StandardScaler()),('m',KNeighborsClassifier(n_neighbors=5))],
}
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
results = []

for name, model in models.items():

    cv_res = cross_validate(
        model,
        X,
        y,
        cv=cv,
        scoring=['accuracy', 'f1_macro', 'roc_auc'],
        return_train_score=False
    )

    results.append({
        'Model': name,
        'Accuracy': cv_res['test_accuracy'].mean(),
        'F1 Score': cv_res['test_f1_macro'].mean(),
        'AUC-ROC': cv_res['test_roc_auc'].mean(),
        'CV Std': cv_res['test_accuracy'].std(),
    })

print(
```

```

        f'{name:22s}: '
        f'Acc={results[-1]["Accuracy"]:.3f} '
        f'F1={results[-1]["F1 Score"]:.3f} '
        f'AUC={results[-1]["AUC-ROC"]:.3f} '

    )

```

```

Logistic Regression : Acc=0.930 F1=0.930 AUC=0.975
Decision Tree       : Acc=0.925 F1=0.925 AUC=0.948
Random Forest       : Acc=0.936 F1=0.936 AUC=0.977
KNN (k=5)           : Acc=0.886 F1=0.886 AUC=0.949

```

```

# ■■■ Build the leaderboard ■■■
df_res = pd.DataFrame(results).sort_values('Accuracy', ascending=False)
df_res[['Accuracy', 'F1 Score', 'AUC-ROC', 'CV Std']] = df_res[['Accuracy', 'F1
print('\n' + '='*70)
print('MODEL LEADERBOARD')
print('='*70)
print(df_res.to_string(index=False))

```

```

=====
MODEL LEADERBOARD
=====

```

	Model	Accuracy	F1 Score	AUC-ROC	CV Std
	Random Forest	0.9362	0.9362	0.9775	0.0310
	Logistic Regression	0.9300	0.9300	0.9752	0.0222
	Decision Tree	0.9250	0.9249	0.9481	0.0265
	KNN (k=5)	0.8863	0.8860	0.9487	0.0242

Hyperparameter Tuning — GridSearchCV

Hyperparameters are settings you choose before training starts — things like `max_depth`, `n_estimators`, `learning_rate`. The model cannot learn these from data; you must set them. `GridSearchCV` automates the search by trying every combination.

Parameters vs Hyperparameters

Feature	Parameters
What are they?	Values learned automatically from the training data
Who decides them?	The machine learning algorithm learns them
Examples	Weights in Logistic Regression, Split Thresholds in Decision Trees, Coefficients
How are they set?	Automatically learned during training
Changed During Training?	Yes, continuously updated while learning
Purpose	Help the model learn patterns from data
Can We Directly Control Them?	Usually No

GridSearchCV — All Combinations

	n=50	n=100	n=200
d=2	0.78	0.82	0.80
d=4	0.84	0.89	0.87
d=6	0.85	0.91 ★ BEST	0.90
d=8	0.84	0.90	0.89

How GridSearchCV Works

1. You give it a grid of values to try.
2. It trains the model for every combination.
3. Each combination is evaluated with K-fold CV.
4. It reports the combination with the best mean CV score.

The table shows all depth × n_estimators combos. **Dark blue = best combo.**

GridSearchCV + RandomizedSearchCV

Grid search tries everything; random search is faster for large grids:

```
# ■ GridSearchCV
from sklearn.model_selection import GridSearchCV, StratifiedKFold
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.datasets import make_classification
import pandas as pd, numpy as np
X, y = make_classification(n_samples=600, n_features=12, random_state=42)
```

```
# ■ Part 1: GridSearchCV for Decision Tree ■
# Hyperparameters to test
param_grid = {
    'max_depth': [2, 4, 6, 8, None],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 3, 5],
    'criterion': ['gini', 'entropy']
}

# 5-Fold Stratified Cross Validation
cv = StratifiedKFold(
    n_splits=5,
    shuffle=True,
    random_state=42
)

# Grid Search
grid_search = GridSearchCV(
    estimator=DecisionTreeClassifier(random_state=42),
    param_grid=param_grid,
    cv=cv,
    scoring='accuracy',
)

# Train and find best combination
grid_search.fit(X, y)

# Results
```



```
print("Best Parameters:", grid_search.best_params_)
print("Best Accuracy:", round(grid_search.best_score_, 4))
```

```
Best Parameters: {'criterion': 'entropy', 'max_depth': 4, 'min_samples_leaf': 1}
Best Accuracy: 0.9083
```

```
# Show GridSearch results
results = pd.DataFrame(grid_search.cv_results_)

# Top 5 models
top_5 = results[
    ['param_max_depth', 'param_criterion',
     'mean_test_score', 'rank_test_score']
].sort_values('rank_test_score').head(5)

print("\nTop 5 Models:")
print(top_5)
```

```
Top 5 Models:
```

	param_max_depth	param_criterion	mean_test_score	rank_test_score
61	4	entropy	0.908333	1
62	4	entropy	0.908333	1
60	4	entropy	0.908333	1
0	2	gini	0.903333	4
4	2	gini	0.903333	4

LAB : MODEL COMPARISION DASHBOARD

Build a complete model selection pipeline: load data, preprocess, train 4 models, evaluate with cross-validation, tune hyperparameters, and produce a final comparison report. We use the Pima Indians Diabetes dataset as the benchmark.

Step 1 — Load, Clean, Split

```
# ■ Load and Clean ■
import pandas as pd, numpy as np, matplotlib.pyplot as plt, seaborn as sns
from sklearn.model_selection import StratifiedKFold, cross_validate, GridSea
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix, Confusi
import time, warnings; warnings.filterwarnings('ignore')

# Check if 'diabetes.csv' exists, if not, download it
import os
if not os.path.exists('diabetes.csv'):
    # Using wget to download the file from a common public repository
```

```

!wget https://raw.githubusercontent.com/plotly/datasets/master/diabetes.

df = pd.read_csv('diabetes.csv')
for col in ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']:
    df[col] = df[col].replace(0, df[col].median())
X = df.drop('Outcome', axis=1)
y = df['Outcome']
print(f'Dataset: {X.shape[0]} rows, {X.shape[1]} features')
print(f'Class balance: {y.value_counts().to_dict()}')
# Output: {0: 500, 1: 268} – slightly imbalanced

--2026-06-05 01:28:22-- https://raw.githubusercontent.com/plotly/datasets/ma
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.10
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.1
HTTP request sent, awaiting response... 200 OK
Length: 23873 (23K) [text/plain]
Saving to: 'diabetes.csv'

diabetes.csv          100%[=====>]  23.31K  --.-KB/s    in 0s

2026-06-05 01:28:22 (81.2 MB/s) - 'diabetes.csv' saved [23873/23873]

Dataset: 768 rows, 8 features
Class balance: {0: 500, 1: 268}

```

Step 2 — Build Models & Run Cross-Validation

```

# ■ Cross-Validation Leaderboard ■
# Models – LR and KNN need scaling; trees don't (Pipeline handles this clean
models = {
    'Logistic Regression': Pipeline([('sc', StandardScaler()),
    ('m', LogisticRegression(max_iter=1000, random_state=42))]),
    'Decision Tree' : DecisionTreeClassifier(max_depth=5, random_state=42),
    'Random Forest' : RandomForestClassifier(n_estimators=100, random_state=42),
    'KNN' : Pipeline([('sc', StandardScaler()),
    ('m', KNeighborsClassifier(n_neighbors=7))]),
}
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
metrics = ['accuracy', 'f1_macro', 'roc_auc', 'precision_macro', 'recall_macro']
records = []
for name, model in models.items():
    t0 = time.time()
    res = cross_validate(model, X, y, cv=cv, scoring=metrics)
    elapsed = time.time()-t0
    row = {'Model': name, 'Time(s)': round(elapsed,2)}
    for m in metrics:
        row[m.replace('_macro', '').title()] = round(res[f'test_{m}'].mean(), 4)
        row[f'{m.replace('_macro', '').title()}_std'] = round(res[f'test_{m}'].std(), 4)
    records.append(row)
leaderboard = pd.DataFrame(records).sort_values('Accuracy', ascending=False)
print(leaderboard[['Model', 'Accuracy', 'F1', 'Roc_Auc', 'Precision', 'Recall', 'T

```

Model	Accuracy	F1	Roc_Auc	Precision	Recall	Time(s)
Logistic Regression	0.7708	0.7320	0.8372	0.7602	0.7219	0.19

Random Forest	0.7577	0.7236	0.8235	0.7367	0.7170	2.78
KNN	0.7435	0.7026	0.7912	0.7206	0.6948	0.14
Decision Tree	0.7082	0.6682	0.7516	0.6851	0.6719	0.12

Step 3 — Tune the Best Model with GridSearchCV

```
# ■ GridSearch Tuning ■
# From the leaderboard, suppose Random Forest wins
# Tune its key hyperparameters
param_grid = {
    'n_estimators' : [50, 100, 200],
    'max_depth' : [None, 5, 10],
    'min_samples_leaf': [1, 3, 5],
}
# 3x3x3 = 27 combos x 5-fold = 135 fits — runs in ~1 minute
gs = GridSearchCV(
    RandomForestClassifier(random_state=42),
    param_grid, cv=cv, scoring='accuracy',
    n_jobs=-1, verbose=0
)
gs.fit(X, y)
print(f'Best params : {gs.best_params_}')
print(f'Best CV acc : {gs.best_score_:.4f}')
# Compare tuned vs default
default_rf = RandomForestClassifier(n_estimators=100, random_state=42)
default_scores = cross_val_score(default_rf, X, y, cv=cv, scoring='accuracy')
tuned_scores = cross_val_score(gs.best_estimator_, X, y, cv=cv, scoring='accuracy')
print(f'\nDefault RF accuracy: {default_scores.mean():.4f} ± {default_scores.std():.4f}')
print(f'Tuned RF accuracy: {tuned_scores.mean():.4f} ± {tuned_scores.std():.4f}')

Best params : {'max_depth': 5, 'min_samples_leaf': 3, 'n_estimators': 50}
Best CV acc : 0.7759

Default RF accuracy: 0.7577 ± 0.0266
Tuned RF accuracy: 0.7759 ± 0.0367
```

Step 4 — Final Evaluation + Confusion Matrix Dashboard


```
# ■ Final Evaluation Dashboard ■
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, ConfusionMatrix

# One final hold-out test — we only do this ONCE
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.20, random_state=42, stratify=y)
fig, axes = plt.subplots(1, 4, figsize=(18, 4.5))
Final_models = {
    'Logistic Regression': Pipeline([('sc', StandardScaler()), ('m', LogisticRegression())]),
    'Decision Tree' : DecisionTreeClassifier(max_depth=5, random_state=42),
    'Random Forest' : gs.best_estimator_, # Tuned version
    'KNN' : Pipeline([('sc', StandardScaler()), ('m', KNeighborsClassifier(n_neighbors=5))])
}
final_results = []
```

```

for ax, (name, model) in zip(axes, Final_models.items()):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    cm = confusion_matrix(y_test, y_pred)
    ConfusionMatrixDisplay(cm, display_labels=['Healthy', 'Diabetic']).plot(
        ax=ax, cmap='Blues', colorbar=False)
    ax.set_title(f'{name}\nAcc: {acc:.2%}', fontsize=9)
    final_results.append({'Model':name, 'Test Accuracy':f'{acc:.2%}'})
plt.suptitle('Final Confusion Matrices – All Models', fontsize=13, y=1.02)
plt.tight_layout(); plt.show()
# Final report for best model
best_model = Final_models['Random Forest']
print('\n=== FINAL REPORT: Random Forest (Tuned) ===')
print(classification_report(y_test, best_model.predict(X_test),
    target_names=['Healthy', 'Diabetic']))

```



```

=== FINAL REPORT: Random Forest (Tuned) ===

```

	precision	recall	f1-score	support
Healthy	0.78	0.85	0.81	100
Diabetic	0.67	0.56	0.61	54
accuracy			0.75	154
macro avg	0.72	0.70	0.71	154
weighted avg	0.74	0.75	0.74	154

Machine Learning Evaluation & Model Selection – Quick Cheat Sheet

Concept	Key Code / Function	
K-Means Clustering	<code>KMeans(n_clusters=K, n_init=10)</code>	Always apply StandardScaler
Elbow Method	<code>km.inertia_</code> for K in <code>range(2,11)</code>	Look for the bend (elbow) w/

Concept	Key Code / Function	
Silhouette Score	<code>silhouette_score(X, labels)</code>	Range: 0 to 1 → Higher = Better
Cluster Visualisation	<code>plt.scatter(X[:,0], X[:,1], c=labels)</code>	Color points by cluster, mark c
Correlation Matrix	<code>df.corr()</code>	Shows feature relationships; c
Correlation Heatmap	<code>sns.heatmap(df.corr(), annot=True)</code>	Strong positive = high positive
Outlier Detection (IQR)	<code>Q1 - 1.5*IQR, Q3 + 1.5*IQR</code>	Detect outliers using lower and
Outlier Capping	<code>df[col].clip(lower, upper)</code>	Winsorizing — reduce outlier i
Outlier Removal	<code>df[(df[col]>=lower)&(df[col]<=upper)]</code>	Remove extreme values if they
Training vs Test Error	<code>model.score(X_train, y_train)</code> <code>model.score(X_test, y_test)</code>	Large gap = Overfitting
Cross Validation	<code>cross_val_score(model, X, y, cv=5)</code>	More reliable than a single tra
Stratified Cross Validation	<code>StratifiedKFold(n_splits=5)</code>	Maintain class balance in clas
Multiple Metrics	<code>cross_validate(model, X, y, scoring=[...])</code>	Evaluate Accuracy, Precision, R
Pipeline	<code>Pipeline([('sc', Scaler()), ('m', Model())])</code>	Prevents data leakage and sim
Model Comparison	Loop + <code>cross_validate()</code> + DataFrame	Compare models using Accur
GridSearchCV	<code>GridSearchCV(model, param_grid, cv=5)</code>	Exhaustive search over all par
RandomizedSearchCV	,	